

Analysis

The result metrics for all 10 search methods are in the table below:

Air Cargo Problem 1						
Search algorithm	Time (sec)	Plan length	Expansions	Goal tests	New nodes	
breadth_first_search	0.2	6	42	57	176	optimal
breadth_first_tree_search	0.54	6	1300	1301	5304	optimal
depth_first_graph_search	0.005	9	10	11	40	
depth_limited_search	0.02	50	52	64	208	
uniform_cost_search	0.2	6	55	57	224	optimal
recursive_best_first_search	1.67	6	4022	4023	16164	optimal
greedy_best_first_graph_search	0.003	6	7	9	28	optimal
A* search with h_1	0.024	6	55	57	224	optimal
A* search with h_ignore_preconditions	0.024	6	41	43	170	optimal
A* search with h_pg_levelsum	0.29	6	11	13	50	optimal
Air Cargo Problem 2						
Search algorithm	Time (sec)	Plan length	Expansions	Goal tests	New nodes	
breadth_first_search	8.7	9	3345	4613	30526	optimal
breadth_first_tree_search						
depth_first_graph_search	0.94	324	329	330	2964	
depth_limited_search	9.48	50	3674	33918	34291	
uniform_cost_search	7.5	9	4780	4782	43381	optimal
recursive_best_first_search						
greedy_best_first_graph_search	0.9	21	598	600	5386	
A* search with h_1	7.5	9	4780	4782	43381	optimal
A* search with h_ignore_preconditions	2.72	9	1450	1452	13303	optimal
A* search with h_pg_levelsum	23.55	9	86	88	841	optimal
Air Cargo Problem 3						
Search algorithm	Time (sec)	Plan length	Expansions	Goal tests	New nodes	
breadth_first_search	67.6	12	14523	17985	127417	optimal
breadth_first_tree_search						
depth_first_graph_search	24.07	2043	5650	5651	45753	
depth_limited_search						
uniform_cost_search	32.82	12	17884	17884	155325	optimal
recursive_best_first_search						
greedy_best_first_graph_search	6.2	21	3373	3375	28947	
A* search with h_1	33.6	12	17882	17884	155325	optimal
A* search with h_ignore_preconditions	10.8	12	5034	5036	44711	optimal
A* search with h_pg_levelsum	135.84	12	315	317	2903	optimal

Optimal plans for **problem 1** were generated by 8 out of 10 methods (please note, not all of them always generate optimal plans). They are different only by the orders of actions. This one was generated by breadth_first_search:

Load(C2, P2, JFK), Load(C1, P1, SFO), Fly(P2, JFK, SFO), Fly(P1, SFO, JFK), Unload(C1, P1, JFK),
Unload(C2, P2, SFO)

5 out of 10 methods returned optimal plans for **problem 2**. The optimal plan generated by A* search with h_ignore_precondition looks as follows:

Load(C1, P1, SFO), Fly(P1, SFO, JFK), Unload(C1, P1, JFK), Load(C2, P2, JFK), Fly(P2, JFK, SFO),
Unload(C2, P2, SFO), Load(C3, P3, ATL), Fly(P3, ATL, SFO), Unload(C3, P3, SFO)

5 out of 10 methods returned optimal plans for **problem 3**. A* search with `h_pg_levelsum` returned the following optimal plan:

Load(C2, P2, JFK), Fly(P2, JFK, ORD), Load(C4, P2, ORD), Fly(P2, ORD, SFO), Load(C1, P1, SFO),
Fly(P1, SFO, ATL), Load(C3, P1, ATL), Fly(P1, ATL, JFK), Unload(C1, P1, JFK), Unload(C2, P2, SFO)
Unload(C3, P1, JFK), Unload(C4, P2, SFO)

Breadth-first search always returns optimal plans and for all three problems it is one of the best search methods in terms of speed and memory. Depth-first search is very fast and has small memory footprint but it produces ridiculously long plans. While for the first problem it might be acceptable (9 steps vs 6 steps in the optimal plan), for the second and third problems the plans are simply useless. Moreover, the lengths of the plans generated by depth-first search are very sensitive to the order of actions in the problem description. The order in my code is (load_actions, fly_actions, unload_actions) and the length of the plan returned by the depth-first search for the third problem is 2043. When I changed the order to the original one: (load_actions, unload_actions, fly_actions) I got a plan with “only” 195 steps. For the second problem, it was 324 vs 1444 respectively.

Another example of the uninformed search algorithms that always produce optimal results is the uniform-cost search algorithm. In our case, for all three problems, the uniform-cost search is, in fact, the breadth-first search because all steps costs are equal to 0. The uniform-cost result metrics are slightly different from the metrics of the breadth-first search because the uniform-cost search uses a priority queue for the frontier vs. FIFO in the breadth-first search.

The other two uninformed search algorithms, breadth-first tree search and depth-limited search, are not very practical for this kind of planning problems. The breadth-first tree search generates an enormous number of duplicated nodes and it never ended for me on the second and third problems. On the first problem, it created 5304 new states vs. 176 states created by the breadth-first search. The depth-limited search never ended on the third problem but returned 50-step plans for the first and second problems (50 was the depth limit). It seems more reasonable to use iterative depth search when we cannot estimate the length of optimal plans

Both informed A* algorithms, A* with `h_ignore_preconditions` and A* with `h_pg_levelsum`, performed very well on all the problems. Both algorithms returned optimal plans. `h_ignore_preconditions` was faster than `h_pg_levelsum` on all 3 problems (0.024 vs 0.29, 2.72 vs 23.55 and 10.8 vs 135.84 respectively) but `h_pg_levelsum` evaluated significantly less nodes (approximately 3.5, 16 and 15.5 times less states for first, second and third problems respectively). The main overhead of the `h_pg_levelsum` algorithm is the planning graph generation. In my implementation, I reduced the generation time by caching positive and negative preconditions and effects that are used in the “A” node mutex functions. We can improve performance even more by ignoring mutexes altogether. Mutexes are used in the GRAPHPLAN algorithm but they are not used in `h_pg_levelsum`. When I disabled `update_a_mutex` and `update_s_mutex` in the `create_graph` method, it took `h_pg_levelsum` only about 40 sec. instead of 135 sec. on the third problem. Another reason why `h_ignore_preconditions` is faster than `h_pg_levelsum` is because in our problems the goals are satisfied by different actions so we can use a very simple and fast implementation (see the code). But in general case (and if we want to get an accurate heuristic), it’s an instance of the set-cover problem which is NP-hard and its greedy polynomial version is not admissible.

In general, the informed search algorithms are superior to the uninformed ones. In our case, the problems are not very complex and so we could use the breadth-first search for all of them. But for harder problems, with hundreds of fluents and actions, the breadth-first search would simply run out of space. `h_ignore_precondition` and (especially) `h_pg_levelsum` expand far less nodes so they should be much more suitable for harder problems.